

# Sliding Context Window

December 15, 2025

# Motivation

- Documentation depends on code, tickets, and discussions; signals are unbounded.

# Motivation

- Documentation depends on code, tickets, and discussions; signals are unbounded.
- LLM context windows (and human attention) are bounded.

# Motivation

- Documentation depends on code, tickets, and discussions; signals are unbounded.
- LLM context windows (and human attention) are bounded.
- Scalability: we must pick the most valuable, self-contained context slice, not dump everything.

# Motivation

- Documentation depends on code, tickets, and discussions; signals are unbounded.
- LLM context windows (and human attention) are bounded.
- Scalability: we must pick the most valuable, self-contained context slice, not dump everything.
- Goal: organize signals to maximize documentation value within hard context limits.

# Scalability $\Rightarrow$ Graph + Optimal Closure

- Chunk the data (messages, issues, PRs, doc gaps) into nodes; remove duplication via dependencies (edges:  $u \rightarrow v$  means  $u$  needs  $v$ ).

## Scalability $\Rightarrow$ Graph + Optimal Closure

- Chunk the data (messages, issues, PRs, doc gaps) into nodes; remove duplication via dependencies (edges:  $u \rightarrow v$  means  $u$  needs  $v$ ).
- Assign utility weight per chunk; higher = more value for the next doc update.

## Scalability $\Rightarrow$ Graph + Optimal Closure

- Chunk the data (messages, issues, PRs, doc gaps) into nodes; remove duplication via dependencies (edges:  $u \rightarrow v$  means  $u$  needs  $v$ ).
- Assign utility weight per chunk; higher = more value for the next doc update.
- Select  $k$  chunks that form a closure (no outgoing edges), maximizing total weight.

# Scalability $\Rightarrow$ Graph + Optimal Closure

- Chunk the data (messages, issues, PRs, doc gaps) into nodes; remove duplication via dependencies (edges:  $u \rightarrow v$  means  $u$  needs  $v$ ).
- Assign utility weight per chunk; higher = more value for the next doc update.
- Select  $k$  chunks that form a closure (no outgoing edges), maximizing total weight.
- This construction **scales** because it enforces bounded chunks, avoids duplication, and optimizes for value under hard context limits.

# Scalability $\Rightarrow$ Graph + Optimal Closure

- Chunk the data (messages, issues, PRs, doc gaps) into nodes; remove duplication via dependencies (edges:  $u \rightarrow v$  means  $u$  needs  $v$ ).
- Assign utility weight per chunk; higher = more value for the next doc update.
- Select  $k$  chunks that form a closure (no outgoing edges), maximizing total weight.
- This construction **scales** because it enforces bounded chunks, avoids duplication, and optimizes for value under hard context limits.
- This is exactly the **maximum-weight closure** problem.

# Optimal Closure Primer

- Given a directed graph with weights, find a closed subset (no outgoing edges) with maximum total weight.

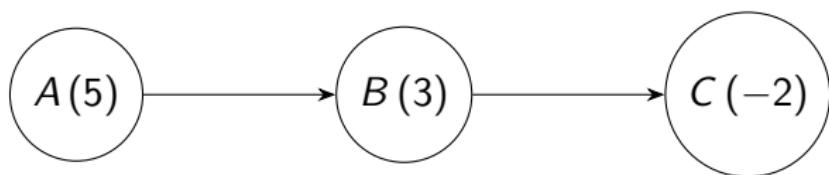
# Optimal Closure Primer

- Given a directed graph with weights, find a closed subset (no outgoing edges) with maximum total weight.
- Solvable via min-cut reduction (polynomial time).

# Optimal Closure Primer

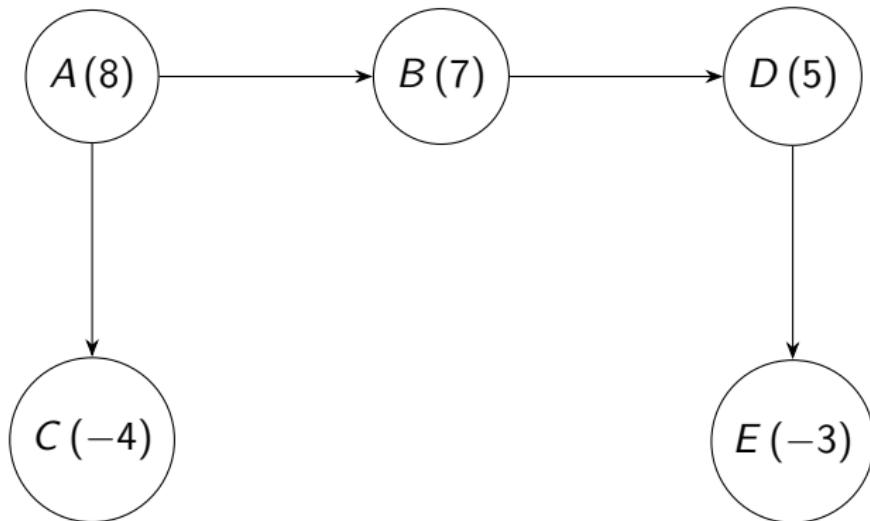
- Given a directed graph with weights, find a closed subset (no outgoing edges) with maximum total weight.
- Solvable via min-cut reduction (polynomial time).
- Our use: pick the best self-contained context under size/budget constraints.

## Example 1 (tiny)



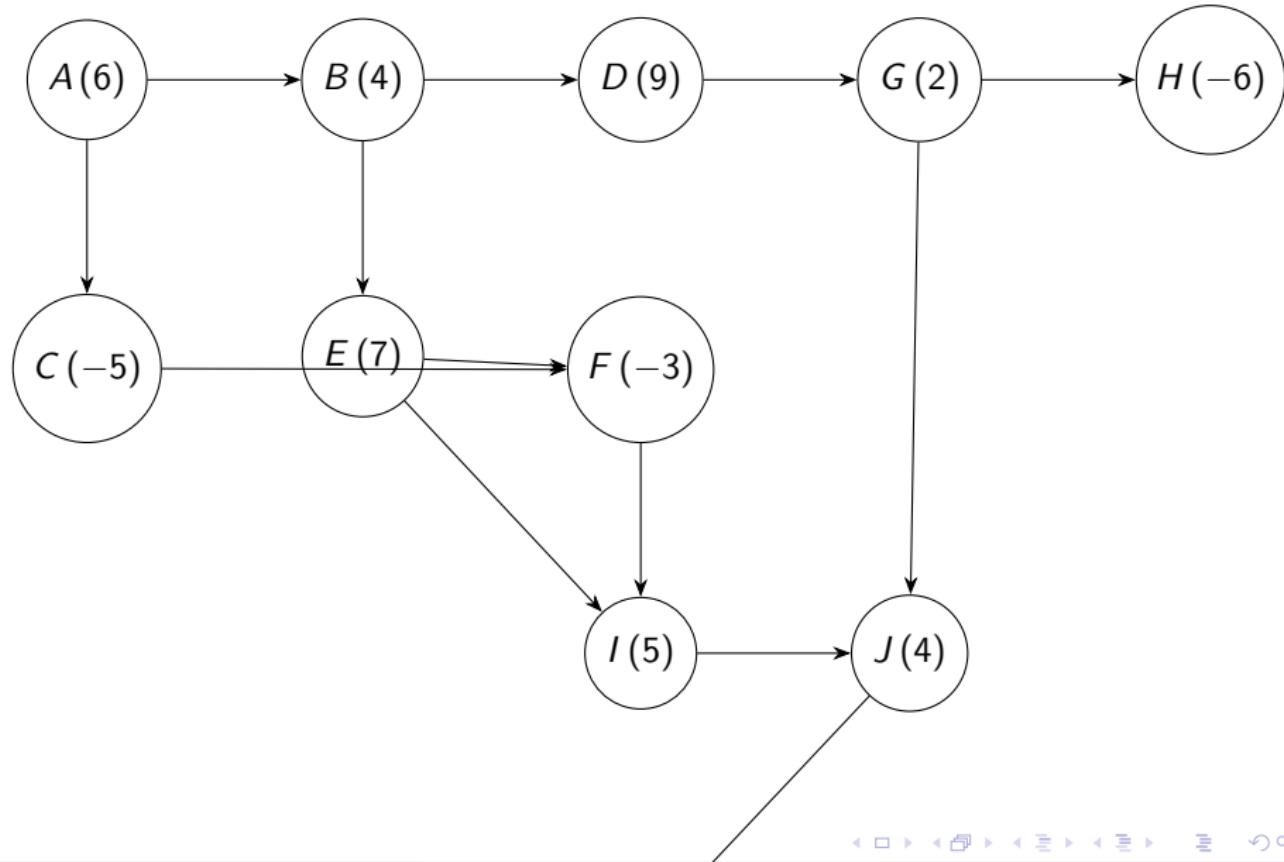
Optimal closure:  $\{A, B, C\}$  (weight 6). This is not hard to inspect.

## Example 2 (medium)



Optimal closure:  $\{A, B, D\}$  (weight 20). This is not hard to inspect, but certainly more complicated than the last one.

Example 3 (larger) ... this problem can be difficult.



# Generalization

- This construction actually solves a more general problem than preparing LLM context windows.
- $k$  need not be an LLM context limit; it can be any resource budget (time, people, memory).

# Generalization

- This construction actually solves a more general problem than preparing LLM context windows.
- $k$  need not be an LLM context limit; it can be any resource budget (time, people, memory).
- Weights can encode any utility (task priority, revenue, risk).

# Generalization

- This construction actually solves a more general problem than preparing LLM context windows.
- $k$  need not be an LLM context limit; it can be any resource budget (time, people, memory).
- Weights can encode any utility (task priority, revenue, risk).
- Closure selection still yields the highest-value, self-contained subset for the chosen budget.

# Generalization

- This construction actually solves a more general problem than preparing LLM context windows.
- $k$  need not be an LLM context limit; it can be any resource budget (time, people, memory).
- Weights can encode any utility (task priority, revenue, risk).
- Closure selection still yields the highest-value, self-contained subset for the chosen budget.
- LLM-driven docs are one special case; the method applies to broader task selection.

# Drawbacks & Mitigations

- Dependency detection may be incomplete or noisy → improve link extraction, allow human corrections.

# Drawbacks & Mitigations

- Dependency detection may be incomplete or noisy → improve link extraction, allow human corrections.
- Super heavy reliance on the Weightings, which may be biased or stale. → add recency decay, source trust, feedback loops.

# Drawbacks & Mitigations

- Dependency detection may be incomplete or noisy → improve link extraction, allow human corrections.
- Super heavy reliance on the Weightings, which may be biased or stale. → add recency decay, source trust, feedback loops.
- Graph churn and scale → incremental updates, caching, and persisted graphs.

# Drawbacks & Mitigations

- Dependency detection may be incomplete or noisy → improve link extraction, allow human corrections.
- Super heavy reliance on the Weightings, which may be biased or stale. → add recency decay, source trust, feedback loops.
- Graph churn and scale → incremental updates, caching, and persisted graphs.
- Empty or sparse edges → fallback visualization only (not for optimization), keep optimization on real edges.

# Drawbacks & Mitigations

- Dependency detection may be incomplete or noisy → improve link extraction, allow human corrections.
- Super heavy reliance on the Weightings, which may be biased or stale. → add recency decay, source trust, feedback loops.
- Graph churn and scale → incremental updates, caching, and persisted graphs.
- Empty or sparse edges → fallback visualization only (not for optimization), keep optimization on real edges.

# System Screenshots I



Figure: Example: simstudioai/sim

*This is a moderately-sized example of a size 24 closure against 120 total nodes.*

# System Screenshots II




Figure: Example: openai/codex

*This is a huge example of a size 500 closure against 923 total nodes.*



# Signal Processing & Storage

- Ingest: GitHub issues/PRs (all open), simple weights (comments + reactions).

# Signal Processing & Storage

- Ingest: GitHub issues/PRs (all open), simple weights (comments + reactions).
- Transform: chunk records with metadata (source, component, tags); edges from dependencies when available.

# Signal Processing & Storage

- Ingest: GitHub issues/PRs (all open), simple weights (comments + reactions).
- Transform: chunk records with metadata (source, component, tags); edges from dependencies when available.
- Store: in-memory per request (current prototype); deterministic layout for graph rendering.

# Signal Processing & Storage

- Ingest: GitHub issues/PRs (all open), simple weights (comments + reactions).
- Transform: chunk records with metadata (source, component, tags); edges from dependencies when available.
- Store: in-memory per request (current prototype); deterministic layout for graph rendering.
- Serve: API returns graph; closure solver returns selected nodes; UI renders graph and selected chunk details.